

A Study of Retrieval-Augmented Generation for Scholarly Question Answering

Ada Lovelace, Grace Hopper

Abstract

This paper investigates retrieval-augmented generation (RAG) as a method for answering research questions grounded in a document corpus. We evaluate hybrid retrieval and reranking strategies.

1. Introduction

Large language models are prone to hallucination when asked about specific facts. Retrieval augmentation mitigates this by grounding generation in retrieved source passages.

2. Methodology

We chunk each paper into paragraph-sized units, embed them with a local sentence-transformers model, and store vectors alongside metadata in an embedded Qdrant collection.

3. Conclusion

Grounding every generated claim in a retrieved chunk is essential for trustworthy scholarly assistants.